

Learning embedded Linux through the file system: A top-down approach

Jason Kridner, jdk@ti.com

June 8, 2011, 10:10AM

ESC Chicago room 34



beagleboard.org

http://beagleboard.org/linux_education

Abstract

- Typical introductions to embedded Linux focus on complex tasks, like building device drivers, that are highly desired by at least some people working on your Linux product team. What about the rest of us that simply want to build a quick understanding of how to do some of the simple tasks that are trivial in microcontrollers without using operating systems, like accessing GPIO pins and performing I2C or SPI serial communications? This class takes a quick tops-down overview from the perspective of an electrical engineer looking to take advantage of the latest in GUI building technologies, like Android, Qt, and HTML5, without completely losing access to the underlying hardware for simple system integration tasks.



Getting started

- Much to learn
 - I'm used to microcontrollers: just give me the datasheet with register definitions and set me free!
- Training on boot & device drivers useful
 - Often geared more at system bring-up
 - What about the everyday user?
 - Where is that abstraction benefit?
- Let's just walk a working system



Kernel.org documentation

<http://www.kernel.org/doc/>

- **Documentation extracted from the Linux kernel and mirrored on the web where Google can find it:**
 - [Documentation](#) - Text files in the kernel source tarball's Documentation subdirectory.
 - [htmldocs](#) - Kernel Documentation maintained in docbook format (output of "make htmldocs").
 - [Menuconfig](#) - help text for each kernel configuration option (from kconfig source).
 - [README](#) various README files scattered around Linux kernel source
 - [RFC](#) - List of IETF RFCs referred to by kernel source files. Links to both the text of the RFC and the source files that refer to it.
 - [Output of kernel's "make help"](#).
- **Standards documents applicable to the Linux kernel**
- **Other web pages containing kernel documentation**
- **Translations to other languages**
- **Documentation on memory management**
- **Miscellaneous**



What is the baseline?

<http://refspecs.linuxfoundation.org/lsb.shtml>

- Every Linux system may be customized
 - This is the nature of open source
 - Stuff still needs to work-together
- The Linux Standard Base
 - Umbrella of various Linux Foundation workgroups
 - A specification and a testkit
 - Documents typical libraries, functions and files expected to be found by the programmer



lsb_release

```
root@beagleboard:~# lsb_release -a
```

```
Distributor ID: Angstrom
```

```
Description:    Angstrom GNU/Linux 2011.03 (Dureza)
```

```
Release:        2011.03
```

```
Codename:       Dureza
```



Filesystem Hierarchy Standard

<http://www.pathname.com/fhs/>

- /bin: essential user command binaries
- /sbin: system binaries
- /boot: static files of the bootloader
- /dev: device files
- /lib: essential shared libraries and kernel modules
- /etc: host-specific configuration
- /home: user home directories
- /root: home directory for the root user
- /opt: add-on application software packages
- /media: mount point for removable media
- /mnt: mount point for temporary mounted file systems
- /tmp: temporary files



Some typical others

- /lost+found: data without directory entry
- /var: data that changes at run-time
- /usr: non-essentials
- /proc: virtual filesystem from kernel
- /sys: virtual filesystem from kernel



/proc

```
root@beagleboard:~# ls /proc
```

1	1151	1214	183	44	87	consoles	kpageflags	sys
10	1160	1215	184	47	9	cpu	loadavg	sysrq-trigger
1021	1171	1218	19	48	901	cpuinfo	locks	sysvipc
1026	1172	1242	2	49	946	crypto	meminfo	timer_list
1029	1177	1251	20	5	950	devices	misc	timer_stats
1031	1182	1254	21	50	953	diskstats	modules	tty
1034	1183	1256	22	51	956	driver	mounts	uptime
1039	1187	1259	23	52	965	execdomains	mtdev	version
1048	1190	1262	24	53	966	fb	net	vmallocinfo
1049	1192	1282	25	57	967	filesystems	pagetypeinfo	vmstat
1061	1193	1295	26	58	993	fs	partitions	zoneinfo
1074	12	1296	27	59	996	interrupts	sched_debug	
1076	1203	13	28	6	asound	iomem	schedstat	
1077	1205	1301	29	61	atags	ioports	scsi	
1081	1207	14	3	62	buddyinfo	irq	self	
1083	1209	15	30	644	bus	kallsyms	slabinfo	
11	1211	16	31	7	cgroups	key-users	softirqs	
1102	1212	17	32	8	cmdline	kmsg	stat	
1121	1213	18	43	857	config.gz	kpagecount	swaps	



/proc

```
root@beagleboard:~# cat /proc/cpuinfo
```

```
Processor           : ARMv7 Processor rev 2 (v7l)
BogoMIPS            : 246.33
Features             : swp half thumb fastmult vfp edsp
                     : thumbee neon vfpv3
CPU implementer     : 0x41
CPU architecture    : 7
CPU variant         : 0x3
CPU part            : 0xc08
CPU revision        : 2

Hardware            : OMAP3 Beagle Board
Revision           : 0020
Serial              : 000000000000000000
```



Kernel Application Binary Interface

<http://www.kernel.org/doc/Documentation/ABI/>

- Low-level kernel interface from “userland”
- Status of interfaces
 - Stable
 - Encouraged to use freely
 - Guaranteed for at least two years
 - Testing
 - Mostly complete, but might change
 - Let developers know how you are using
 - Obsolete
 - Scheduled for removal
 - Removed



Kernel Application Binary Interface

<http://www.kernel.org/doc/Documentation/ABI/>

- Types of interfaces
 - Syscalls
 - Trap interface with IDs
 - May be possible to have a direct entry
 - SYSFS
 - Virtual file system



The file interface abstraction

- What can I do with files?
 - open, read, write, close, delete
 - What is an 'ioctl'?
- What is a virtual file system?
 - Looks like a file, but executes code in the driver
 - Not really storing anything to media
 - A bit like a “ram disk”
 - What is 'mmap'?



www.shutterstock.com · 59109898



Busybox: The Swiss Army Knife of Embedded Linux

<http://www.busybox.net/>

- Multiple execution binary using symbolic links

- Examples:

- [, [[, addgroup, adduser, ar, ash, awk, basename, blkid, bunzip2, bzip2, cat, chattr, chgrp, chmod, chown, chpasswd, chroot, chvt, clear, cmp, cp, cpio, cryptpw, cut, date, dc, dd, deallocvt, delgroup, deluser, df, dhcprelay, diff, dirname, dmesg, du, dumpkmap, dumpleases, echo, egrep, env, expr, false, fbset, fbsplash, fdisk, fgrep, find, free, freeramdisk, fsck, fsck.minix, fuser, getopt, getty, grep, gunzip, gzip, halt, head, hexdump, hostname, httpd, hwclock, id, ifconfig, ifdown, ifup, init, insmod, ip, kill, killall, klogd, last, less, linuxrc, ln, loadfont, loadkmap, logger, login, logname, logread, losetup, ls, lsmod, makedevs, md5sum, mdev, microcom, mkdir, mkfifo, mkfs.minix, mknod, mkswap, mktemp, modprobe, more, mount, mv, nc, netstat, nice, nohup, nslookup, od, openvt, passwd, patch, pidof, ping, ping6, pivot_root, poweroff, printf, ps, pwd, rdate, rdev, readahead, readlink, readprofile, realpath, reboot, renice, reset, rm, rmdir, rmmmod, route, rtcwake, run-parts, sed, seq, setconsole, setfont, sh, showkey, sleep, sort, start-stop-daemon, stat, strings, stty, su, sulogin, swapoff, swapon, switch_root, sync, sysctl, syslogd, tail, tar, tee, telnet, telnetd, test, tftp, time, top, touch, tr, traceroute, true, tty, udhcpc, udhcpd, umount, uname, uniq, unzip, uptime, usleep, vi, vlock, watch, wc, wget, which, who, whoami, xargs, yes, zcat



Syscalls

<http://www.kernel.org/doc/man-pages/online/pages/man2/syscalls.2.html>

- open/read/write/lseek/unlink
- ioctl
- mknod
- fork/select/poll/...
- mkdir/...



What is SYSFS?



www.shutterstock.com · 43373236

- Virtual file system that exposes drivers to userspace
- mount
 - sysfs on /sys type sysfs (rw,relatime)
- /sys/devices ← driver hierarchy
- /sys/bus ← symbolic links to bus owners
- /sys/class ← common interfaces
- /sys/block ← block interface
- How about some examples...



/sys/module

<http://www.kernel.org/doc/Documentation/ABI/stable/sysfs-module>

- /sys/module/MODULENAME
 - .../parameters: options you can provide
 - .../refcnt: number of times in use
- Examples:

8250	input_polldev	omapdss	snd	ubi
auth_rpcgss	ipv6	omapfb	snd_pcm	ubifs
block	kernel	oprofile	snd_pcm_oss	uinput
bluetooth	keyboard	printk	snd_rawmidi	usb_storage
btusb	lockd	psmouse	snd_timer	usbcore
dns_resolver	mmc_core	rfcomm	snd_usb_audio	usbhid
eeprom_93cx6	mmcblk	rfskill	soundcore	videobuf_core
ehci_hcd	mousedev	scsi_mod	spidev	vpfe_capture
hci_uart	musb_hsrc	sg	spurious	vt
hid	netpoll	smc911x	sr_mod	xz_dec
hid_apple	nfs	smc91x	sunrpc	
hid_ntrig	omap3_isp	smc911x	tcp_cubic	
hidp	omap_vout	smc95xx	uas	



Discovering USBFS...

```
root@beagleboard:~# cat /sys/module/usbcore/parameters/usbfs_snoop
N
root@beagleboard:~# echo Y > /sys/module/usbcore/parameters/usbfs_snoop
root@beagleboard:~# mount -t usbfs none /mnt
root@beagleboard:~# ls /mnt
001 Devices
root@beagleboard:~/labs# cat /mnt/devices
```

```
T: Bus=01 Lev=00 Prnt=00 Port=00 Cnt=00 Dev#= 1 Spd=480 MxCh= 3
B: Alloc= 0/800 us ( 0%), #Int= 3, #Iso= 0
D: Ver= 2.00 Cls=09(hub ) Sub=00 Prot=00 MxPS=64 #Cfgs= 1
P: Vendor=1d6b ProdID=0002 Rev= 2.06
S: Manufacturer=Linux 2.6.39 ehci_hcd
S: Product=OMAP-EHCI Host Controller
S: SerialNumber=ehci-omap.0
C:* #Ifs= 1 Cfg#= 1 Atr=e0 MxPwr= 0mA
I:* If#= 0 Alt= 0 #EPs= 1 Cls=09(hub ) Sub=00 Prot=00 Driver=hub
E: Ad=81(I) Atr=03(Int.) MxPS= 4 Iv1=256ms
```

```
T: Bus=01 Lev=02 Prnt=02 Port=03 Cnt=03 Dev#= 5 Spd=1.5 MxCh= 0
D: Ver= 1.10 Cls=00(>ifc ) Sub=00 Prot=00 MxPS= 8 #Cfgs= 1
P: Vendor=093a ProdID=2510 Rev= 1.00
S: Manufacturer=PixArt
S: Product=USB Optical Mouse
C:* #Ifs= 1 Cfg#= 1 Atr=a0 MxPwr=100mA
I:* If#= 0 Alt= 0 #EPs= 1 Cls=03(HID ) Sub=01 Prot=02 Driver=usbhid
E: Ad=81(I) Atr=03(Int.) MxPS= 4 Iv1=10ms
```



/sys/class/leds

<http://www.kernel.org/doc/Documentation/ABI/testing/sysfs-class-led>

- /sys/class/leds/LED
 - .../brightness: 0-max_brightness, >0 = on
 - .../max_brightness: default is 255
 - .../trigger: triggers available from kernel
 - .../inverted: invert on/off state

```
root@beagleboard:~# cat /sys/class/leds/beagleboard\:\:<TAB><TAB>
/sys/class/leds/beagleboard::pmu_stat/ /sys/class/leds/beagleboard::usr0/
/sys/class/leds/beagleboard::usr1/
root@beagleboard:~# cat /sys/class/leds/beagleboard\:\:usr0/trigger
none nand-disk mmc0 [heartbeat]
root@beagleboard:~# echo "none" > /sys/class/leds/beagleboard\:\:usr0/trigger
root@beagleboard:~# echo 1 > /sys/class/leds/beagleboard\:\:usr0/brightness
root@beagleboard:~# cat /sys/class/leds/beagleboard\:\:usr0/brightness
1
root@beagleboard:~# echo "heartbeat" > /sys/class/leds/beagleboard\:\:usr0/trigger
```



beagleboard.org



www.shutterstock.com · 45826603

/sys/class/gpio

<http://www.kernel.org/doc/Documentation/ABI/testing/sysfs-gpio>

- Must be explicitly exported to userspace and not claimed by kernel code
- /sys/class/gpio
 - .../export: asks the kernel to export a GPIO to userspace
 - .../unexport: to return a GPIO to the kernel
 - .../gpioN: for each exported GPIO #N
 - .../value: always readable, writes fail for input GPIOs
 - .../direction: r/w as: in, out (low); write: high, low
 - .../edge: r/w as: none, falling, rising, both
 - .../gpiochipN: for each gpiochip; #N is its first GPIO
 - .../base: (r/o) same as N
 - .../label: (r/o) descriptive, not necessarily unique
 - .../ngpio: (r/o) number of GPIOs; numbered N to N + (ngpio - 1)



GPIOs

readgpio

Board rev	I D	GPIO 173	GPIO 172	GPIO 171
Ax/Bx	7	1	1	1
C1/2/3	6	1	1	0
C4	5	1	0	1
xM-Ax	0	0	0	0

- `cd /sys/class/gpio; ls`
- `echo "171" > export; ls`
- `echo "in" > gpio171/direction`
- `cat gpio171/value`
- `readgpio 172**`

** ^C to cancel



GPIOs

readgpio

- `cd /sys/class/gpio; ls`
- `echo "7"* > export**; ls`
- `echo "in" > gpio7/direction`
- `cat gpio7/value`

• * on xM, use GPIO 4

• ** won't work if kernel has allocated the GPIO already as an event, so you'd need an older kernel to do this (or just point to an expansion header pin)



Reading events

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char** argv)
{
    int fd, ret, data[64];
    fd = ret = open(argv[argc - 1], O_RDONLY);
    printf("open returned %d\n", ret);
    while (1) {
        ret = read(fd, &data, sizeof(data));
        printf("read returned %d\n", ret);
        sleep(1);
    }

    exit(0);
}
```



Reading events

```
root@beagleboard:~/labs# gcc -o myread myread.c
root@beagleboard:~/labs# ./myread /dev/input/event0
open returned 3
read returned 32
read returned 32
read returned 32
root@beagleboard:~/labs# evtest /dev/input/event0
Input driver version is 1.0.1
Input device ID: bus 0x19 vendor 0x1 product 0x1 version 0x100
Input device name: "gpio-keys"
Supported events:
    Event type 0 (Sync)
    Event type 1 (Key)
        Event code 276 (ExtraBtn)
Testing ... (interrupt to exit)
Event: time 1307366684.251618, type 1 (Key), code 276 (ExtraBtn), value 1
Event: time 1307366684.251649, ----- Report Sync -----
Event: time 1307366684.348203, type 1 (Key), code 276 (ExtraBtn), value 0
Event: time 1307366684.348234, ----- Report Sync -----
```




Read events

testuserbtn

- Kernel documentation:

<http://git.kernel.org/?p=linux/kernel/git/torvalds/linux-2.6.git;a=blob;f=Documentation/input/input.txt>


- opkg install evtest
- evtest /dev/input/event0
 - Press the “USER” button
 - ^C to exit
- evtest /dev/input/event2
 - Move the mouse
 - ^C to exit



Access monitor EDID EEPROM

testedid

- `cd /sys/bus; ls`
- `cd i2c/devices; ls`
- `echo "eeprom 0x50" > i2c-3/new_device; ls`
- `cat 3-0050/eeprom**`
- `opkg install i2c-tools`
- `i2cdump -y 0x3 0x50 b`
- `decode-edid**`
- `fbset`



** requires newer kernel than provided



USB OTG and EHCI

- `cd /sys/bus/usb/devices; ls`
- `cat usb1/speed`
- `cat usb?/manufacturer`
- `lsusb`



Natural User Interface



Device nodes

- Kernel interfaces that can be exposed to a file handle
 - Instead of hierarchical directory structure, uses a set of major and minor numbers
- Google udev
- `mknod --help`



Is everything really just a file?

- No, but it is the primary model
- The kernel also provides functions through a vector table used by the C library
- If you create a new driver, you should create SYSFS entries to enable easy use




www.shutterstock.com · 51962428



Lab #1 – Toggle an LED with SYSFS

- Now it is your turn to play...
- `cd lab1; less README.txt`
- 'q' to exit less
- Perform the instructions



GUI building with Qt

Qt C++ framework is just one option for creating graphical applications, but it is fast, flexible, cross-platform and well-supported by an open source community



beagleboard.org

Thanks to Gregg Lebovitz of ICS

Qt architecture



Qt modular class library

Core	XML
GUI	Multimedia
WebKit	Database
Graphics View	Network
Scripting	Unit Tests
OpenGL	Benchmarking

Qt development tools



Qt Creator
Cross-platform IDE



Qt Designer
GUI designer



Qt Assistant
Help reader



Qt Linguist
I18N Toolset

qmake
Cross-Platform
Build Tool

Cross-platform support

Windows

Mac

Linux/X11

Embedded
Linux

Win CE

S60*

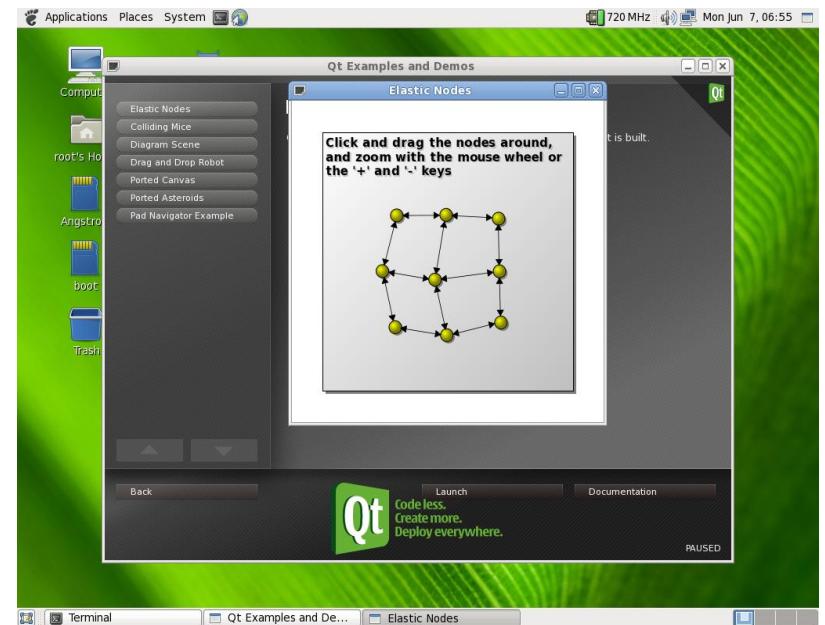
**Coming soon*

Chipsets



Qt demos

- qtdemo
- Descriptions not compiled in here
- Lots of different GUI tools
- Many other programming tools, like networking, IPC, 3D, database, ...



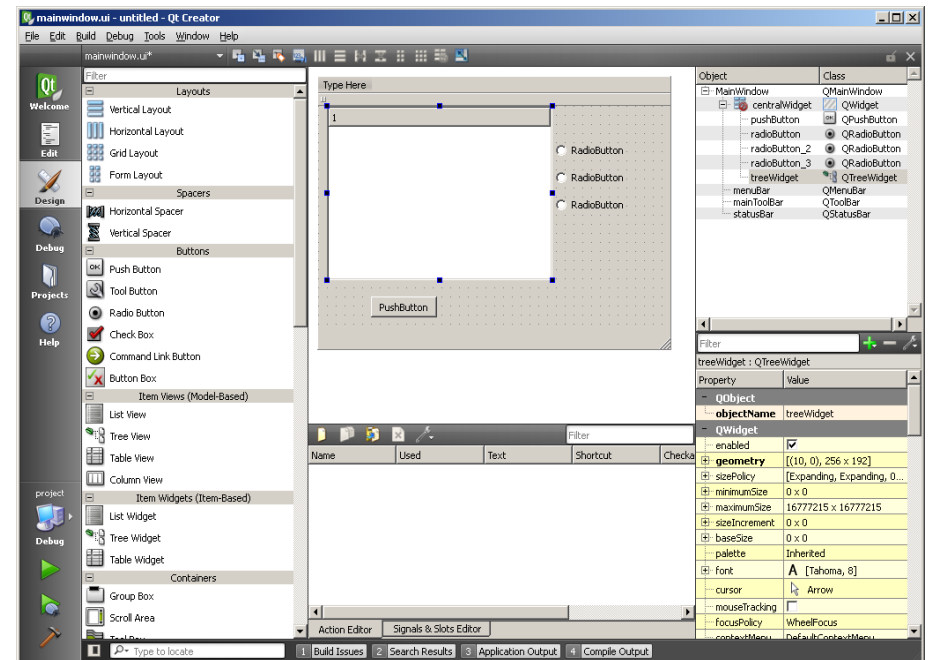
Qt Creator

<http://qt.nokia.com>

- Integrated development environment
 - Runs on Windows, Mac, or Linux
 - Designer for your GUIs
 - C++ editor and debugger
- Build your GUI on your PC, then move it over to the BeagleBoard to add I/O, etc.
 - Angstrom Linux distribution has the compiler and libraries ready to build Qt apps natively on your BeagleBoard



beagleboard.org



The Hello World example


```
#include <QtGui>

int main( int argc, char *argv[] )
{
    QApplication myapp( argc, argv );

    QLabel *mylabel = new
    QLabel("Helloworld");

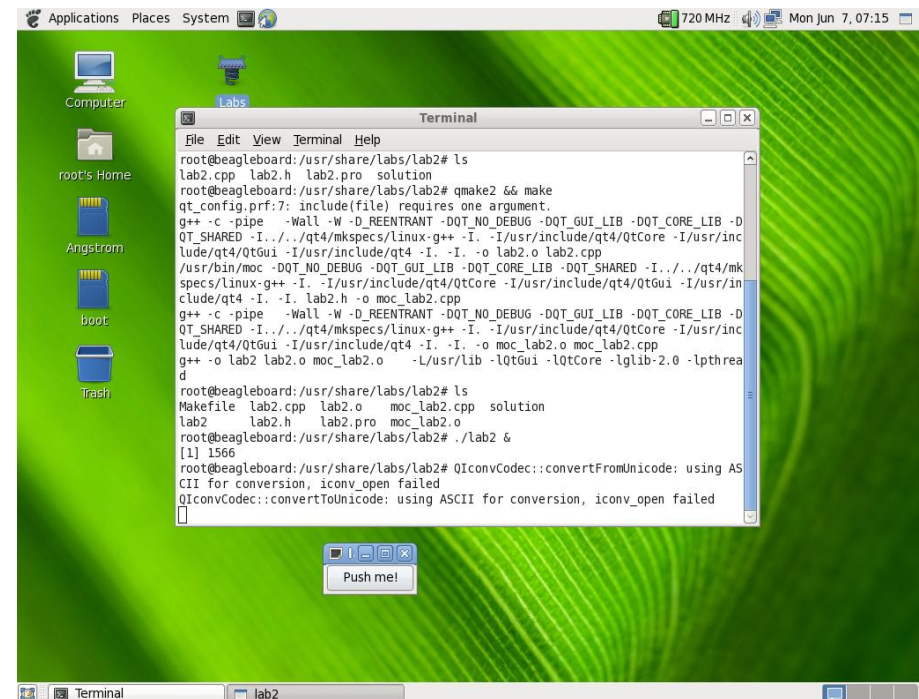
    mylabel->show();

    return myapp.exec();
}
```



Lab #2 – Toggle LED from a GUI

- Close and re-open our ‘Labs’ terminal
- `cd lab2; ls`
- `qmake2 && make`
- `./lab2`
- `gedit ./lab2.cpp &`
- Edit `lab2.cpp` to turn the heartbeat on and off
- `make`
- `./lab2`



Summary

- Utilize <http://www.kernel.org/doc>
- Google!

